

Serverless ML Deployment on AWS

Šimon Soták

- Studied Computer Graphics at Charles University in Prague
- Started out in Game Development (Keen SWH, Warhorse)
- Joined Represent in 2014



Andrej Hoos

- Studied Computer Science and Physics at University of Glasgow
- Joined Represent in 2014





Represent

Arnold Schwarzenegger

\$1M

revenue

69K

products sold





CustomInk kupuje za několik miliard
korun startup se slovenským
spoluzakladatelem



PRG

Prague, Czech Republic
Šaldova 12, Karlín



LA

Los Angeles, California
1680 Vine St, Hollywood

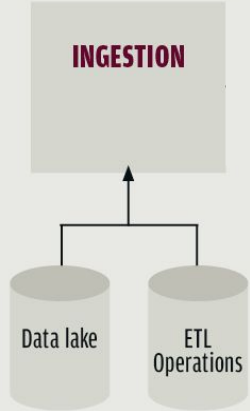


DC

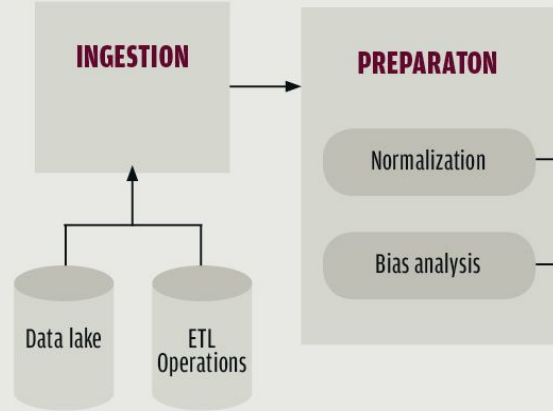
Washington metropolitan area
2910 District Ave, Fairfax

Machine Learning

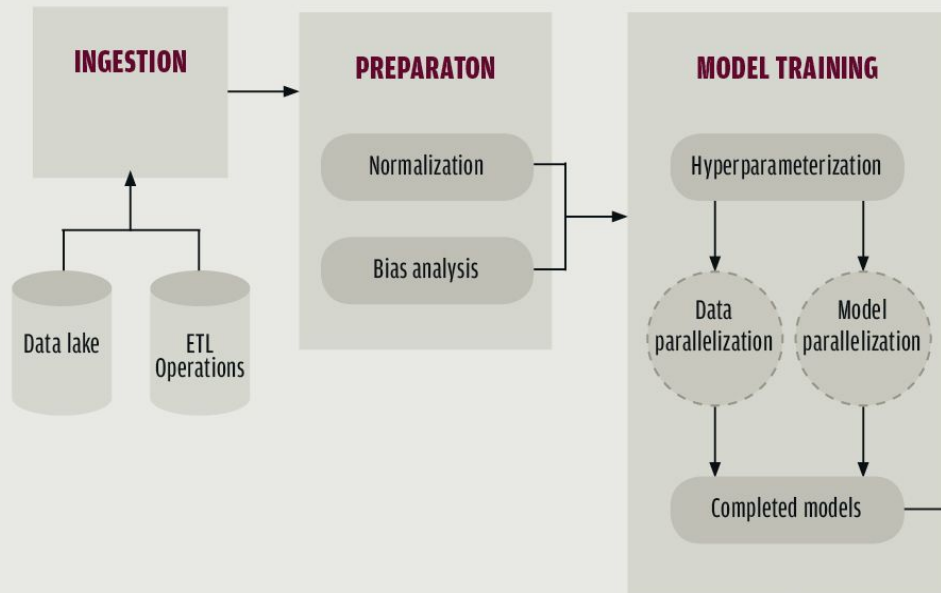
Machine Learning Process



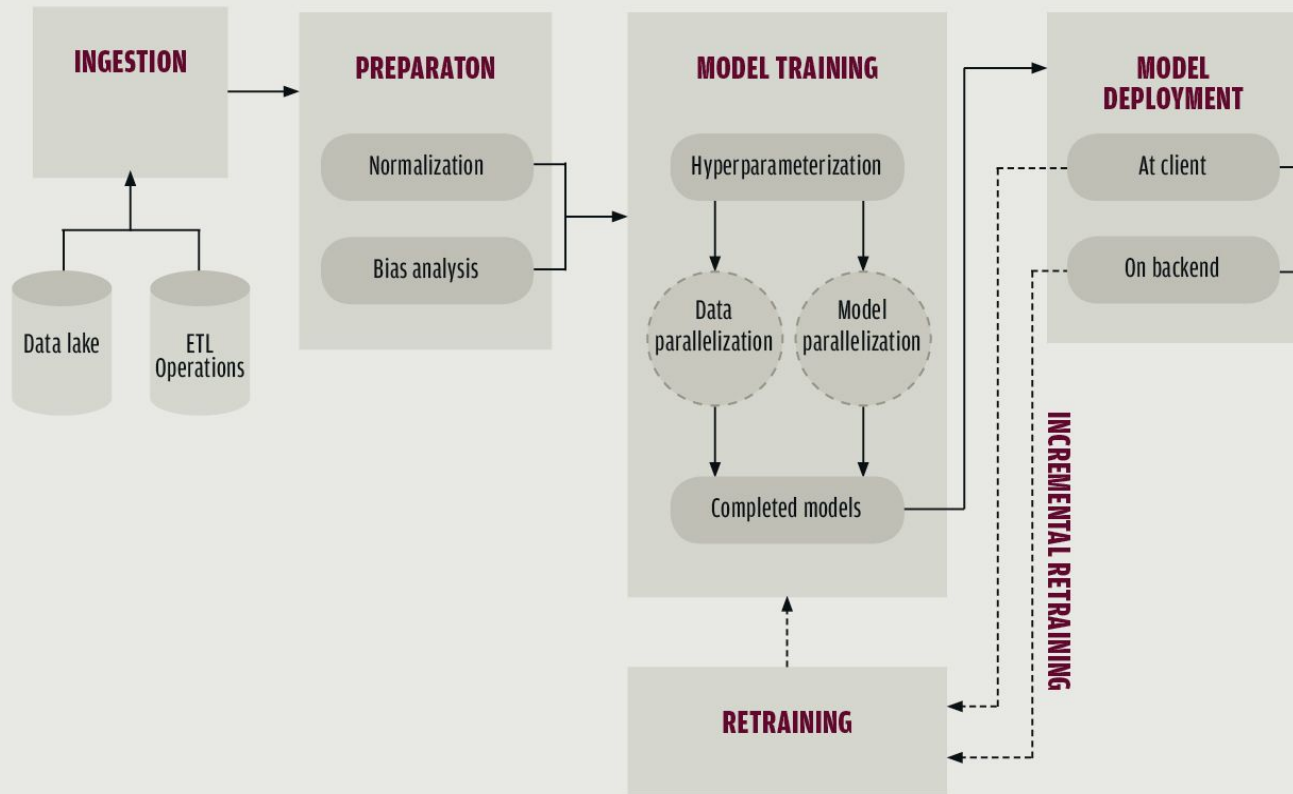
Machine Learning Process



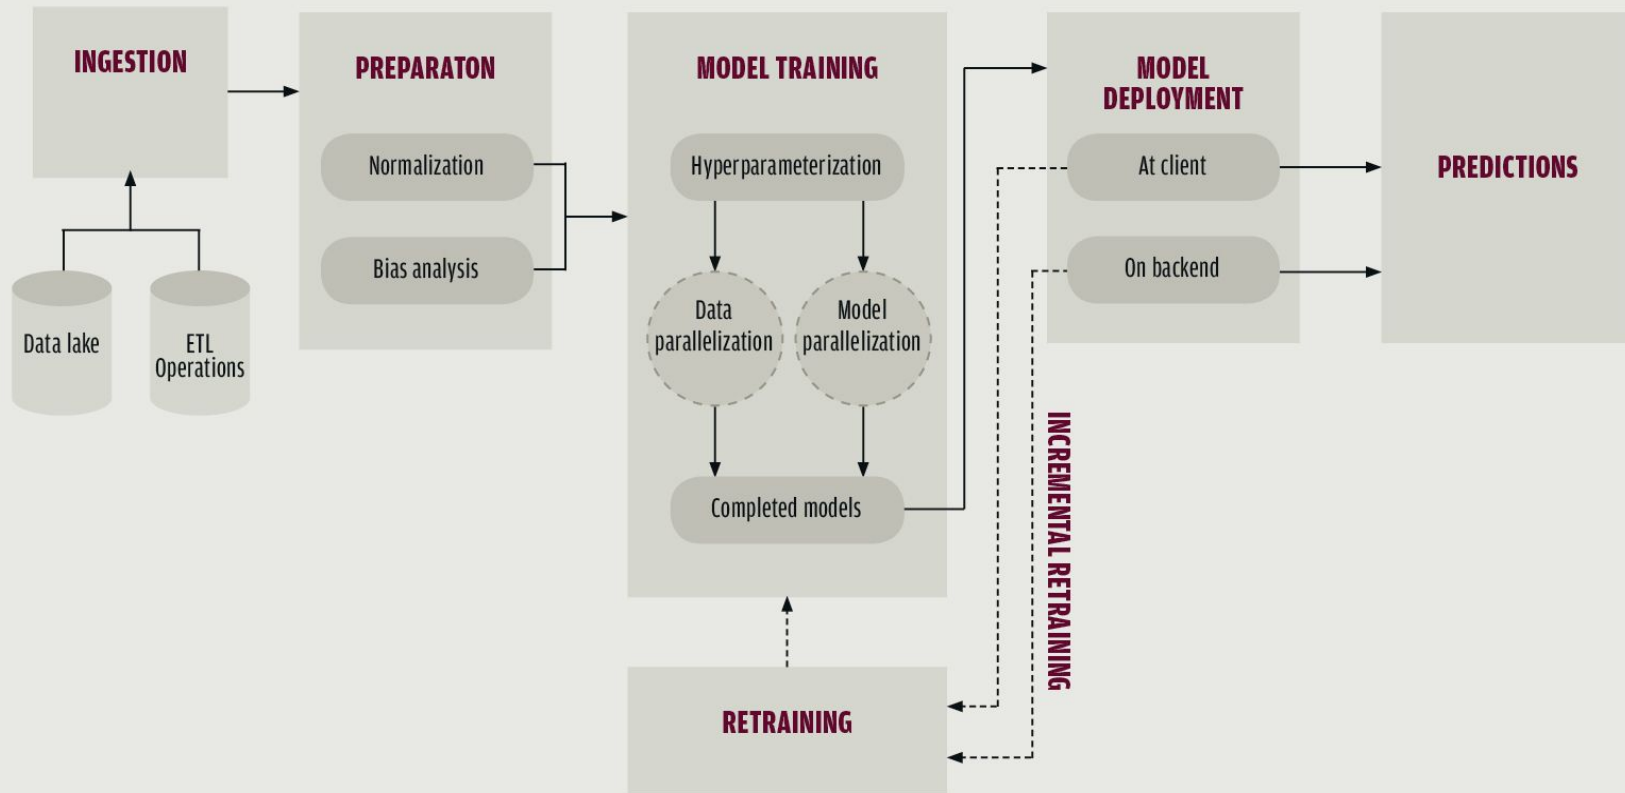
Machine Learning Process

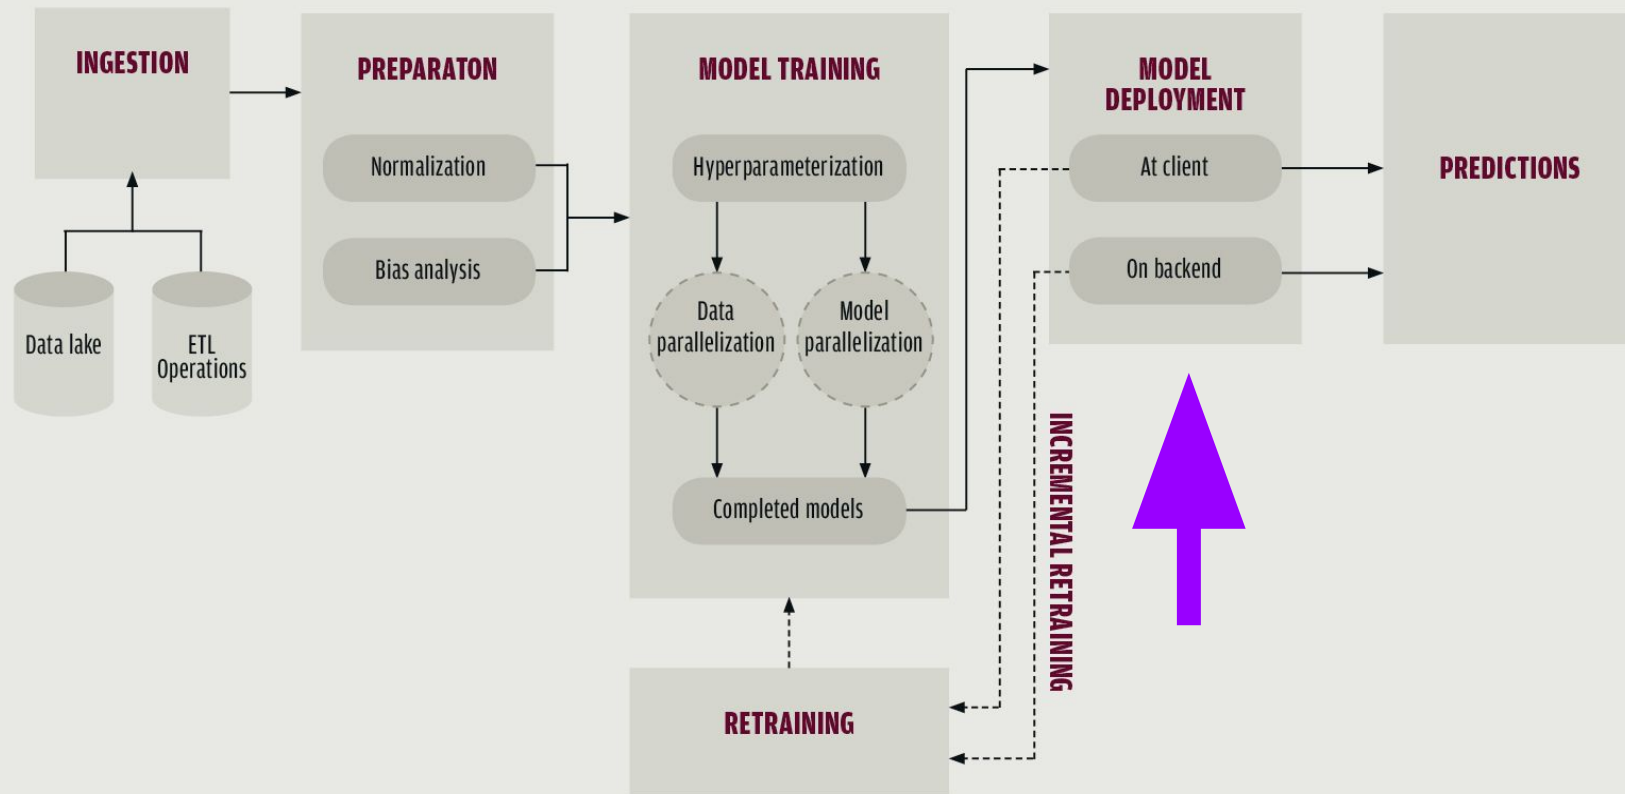


Machine Learning Process



Machine Learning Process





Serverless ML Deployment on AWS

How to deploy ML models

Approach	AWS Service
Hosted servers	EC2
Docker / Kubernetes containers	ECS / EKS
Serverless	Lambda / Fargate
ML-tailored solutions	Sagemaker

What is serverless?

- Function as a Service
- Fully managed
- You provide only the code to be run
- Advantages
 - No maintenance required
 - Scaling “for free”
 - Easier development
 - Lower infrastructure cost

Python ML Deployment on Serverless

- AWS Lambda
- Why Amazon?
- Why Python?

So, can I host my Python ML
model on AWS Lambda?

Yes*...

Cold Start

- First request – “bootup” & init
- Instance “shuts down” after a while if no requests arrive
- Parallel cold starts
- Pre-warming
- ML Model loading

Issue to overcome: Package size

- AWS Lambda – 250 MB max package size
- Scikit-learn – 200+ MB
- Tensorflow – over 250 MB

Did you say
“Easier development”
earlier?

Introducing: `lambdipy`

- Tool for building python packages for AWS Lambda
- github.com/customink/lambdipy
- Builds “minified” packages for AWS Lambda
- Contains pre-built popular packages
- Active development – contributions welcome

What does `lambdipy` do

- Builds any python ML package
 - Uses Docker with Lambda-like environment
 - Downloads pre-built popular packages
- **strip** binaries
 - Binaries take the most space in ML packages
 - This reduces their file size significantly
- Remove some unneeded files
 - Tests, caches, package metadata
- `scipy`: 140 MB → 33 MB

Example

MNIST dataset

label = 5



label = 0



label = 4



label = 1



label = 9



label = 2



label = 1



label = 3



label = 1



label = 4



label = 3



label = 5



label = 3



label = 6



label = 1



label = 7



label = 2



label = 8



label = 6



label = 9



Simple image classifier

- Handwritten digit classification
- Inspired by blog post by Bikramjot Singh Hanzra¹
- **sam** for deployment
- **pipenv** for dependency management

1. <http://hanzratech.in/2015/02/24/handwritten-digit-recognition-using-opencv-sklearn-and-python.html>

Example: train.py

```
dataset = datasets.fetch_mldata("MNIST Original")
```

Example: train.py

```
dataset = datasets.fetch_mldata("MNIST Original")  
  
# Train a ML model  
model = LinearClassifier()  
model.fit(features, labels)
```

Example: train.py

```
dataset = datasets.fetch_mldata("MNIST Original")

# Train a ML model
model = LinearClassifier()
model.fit(features, labels)

joblib.dump(model, "model.pkl")
```

Example: predict.py

```
model = joblib.load("model.pkl")
```


Example: predict.py

```
model = joblib.load("model.pkl")
```

```
def predict(context, _):
```

Example: predict.py

```
model = joblib.load("model.pkl")
```

```
def predict(context, _):  
    request_body = json.loads(context["body"])
```

Example: predict.py

```
model = joblib.load("model.pkl")
```

```
def predict(context, _):  
    request_body = json.loads(context["body"])  
  
    # Load image from url and extract its features  
    prediction = model.predict(image_features)
```

Example: predict.py

```
model = joblib.load("model.pkl")

def predict(context, _):
    request_body = json.loads(context["body"])

    # Load image from url and extract its features
    prediction = model.predict(image_features)

    return {
        "statusCode": 200,
        "body": json.dumps({"digit": int(prediction[0])})
    }
```

Example: template.yml

```
...

Resources:
  PredictFunction:
    Type: AWS::Serverless::Function
    Properties:
      CodeUri: ./build
      Handler: predict.predict
```

...

Example: deployment

- Install dependencies
 - `pipenv install scikit-learn scikit-image numpy`
 - `pipenv install lambdipy aws-sam-cli --dev`
- Train the classifier
 - `pipenv run python train.py`
- Build the bundle
 - `pipenv run lambdipy build --from-pipenv -i predict.py -i model.pkl`
- Deploy
 - `pipenv run sam package --s3-bucket ...`
 - `pipenv run sam deploy --stack-name ...`



Live Demo



Example: testing the predictions

POST

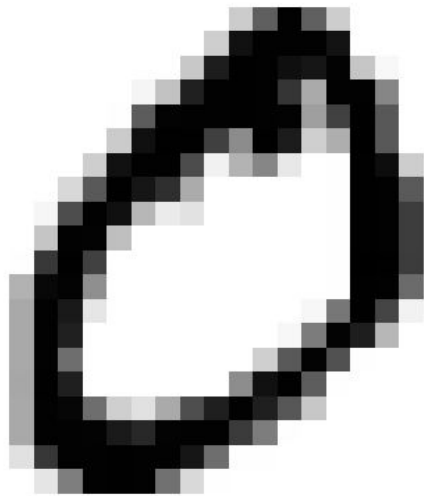
<https://r7my4vk3t0.execute-api.us-east-1.amazonaws.com/Prod/predict>

```
{  
  "url": "https://i.stack.imgur.com/FK0FB.png"  
}
```



```
{ "digit": 0 }
```

```
curl -d '{"url": "https://i.stack.imgur.com/FK0FB.png"}' -X POST  
https://r7my4vk3t0.execute-api.us-east-1.amazonaws.com/Prod/predict
```



Example: testing the predictions

POST

<https://r7my4vk3t0.execute-api.us-east-1.amazonaws.com/Prod/predict>

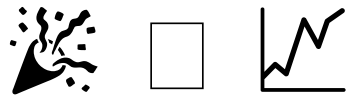
```
{  
  "url": "https://goo.gl/K9ditV"  
}
```



```
{ "digit": 2 }
```

```
curl -d '{"url": "https://goo.gl/K9ditV"}' -X POST  
https://r7my4vk3t0.execute-api.us-east-1.amazonaws.com/Prod/predict
```





Production ML system complete!

Alternative approaches

- Zappa
 - Deploys Python web apps to serverless environments
- Google Cloud Functions
 - Python support currently in Beta

Thank you!

This presentation: tiny.cc/aws-lambdipy

We are hiring and always looking to have good conversations

represent.com/join

